

# DATASCI 207

Nedelina Teneva, PhD  
[nteneva@berkeley.edu](mailto:nteneva@berkeley.edu)

School of Information, UC Berkeley

# Announcements

- Finalize datasets by next week: enter dataset info in the [Logistics Sheet](#)
- No class on July 4th

# HW Recap

- HW3
  - Q3: Data normalization: not **using the train mean/std dev to normalize test data**
  - Q4: Report **validation** loss for learning rates
  - Q5: RMSE is wrong due to incorrect normalization
- HW4
  - Q3: forgot to run the loss on train/test

# Generative vs Discriminative Models

[On Discriminative vs. Generative Classifiers: A comparison of Logistic Regression and Naive Bayes](#), Ng & Jordan



- **Generative classifiers**
  - learn a model of the joint probability  $p(x,y)$  and pick the most likely  $y$
  - Make predictions using Bayes rule to compute the posterior  $p(y|x)$  and from it choose the most likely  $y$ .
  - Examples: Naive Bayes (earlier in the semester), Gaussian Mixture Models (later on)
- **Discriminative classifiers:**
  - Model the posterior  $p(y|x)$  directly and learn a map from the inputs ( $x$ ) to the labels ( $y$ )
  - Examples: logistic regression, trees, Support Vector Machines (SVM)

Generative models need to compute  $p(x|y)$  for Bayes theorem to work – this is an extra step, which discriminative classifiers avoid by modeling  $p(y|x)$  directly (Bayesian <> Frequentists debate)

# Ensembles

# Learning with Ensembles


- The idea is to combine different classifiers into a super-classifier whose performance is much better than using each individual classifier alone.
- Assume that you collect predictions from 10 expert classifiers (e.g., KNN, NB, ...). ensemble methods will allow us to combine their individual predictions to come up with a final prediction that is more accurate and robust.
- 3 most commonly used ensembles: majority vote, bagging, and boosting.
  - **Bagging**: train models in parallel via bootstrap sampling
  - **Boosting**: train additive models in series where each predicts the residual from the previous ones



# Majority Vote

The majority vote principle

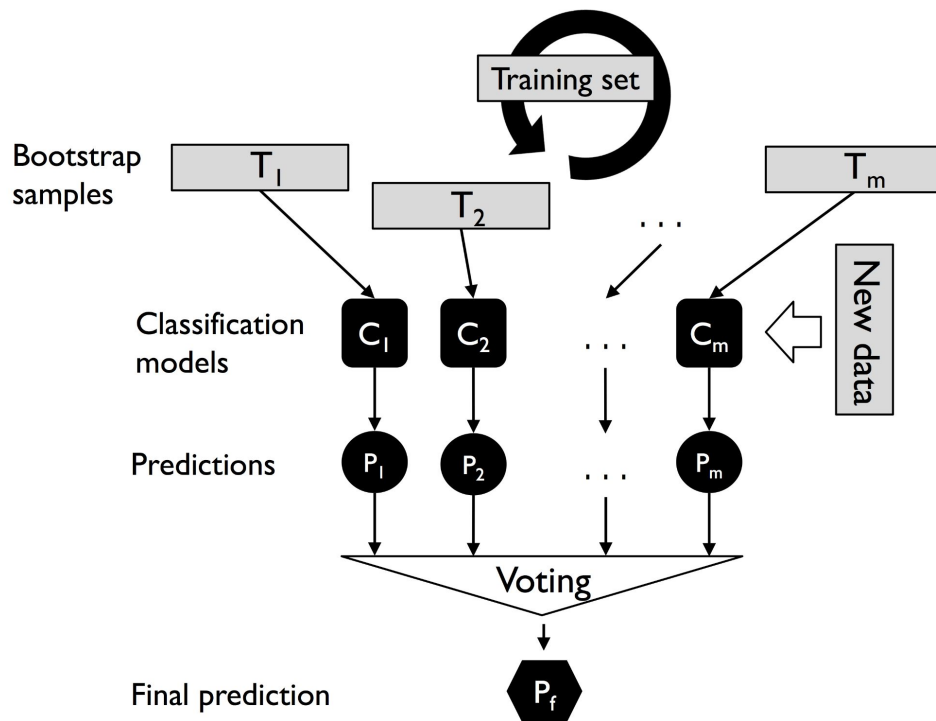
- selects the class label that has been predicted by the majority of classifiers (i.e., received more than 50 percent of votes).
- the term majority vote is used for a binary class setting.
- plurality vote is used for a multiclass setting.



# Bagging

## The bagging principle

- is closely related to the majority vote technique.
- the difference is that instead of using the same training data to fit the different classifiers in the ensemble, **we draw bootstrap (random) samples with replacement** from the initial training data.





# Adaboost

## The adaptive boosting principle - AdaBoost classifiers

- the ensemble consists of very simple base classifiers (**weak** learners; think decision tree stump).
- focuses on training examples that are hard to classify, i.e. let the weak learners learn from their mistakes (misclassified training examples) to improve the performance of the ensemble.
- in contrast to bagging, the boosting algorithm uses random subsets of training examples drawn from the training dataset **without** replacement.

example decision stump: tree of height 1

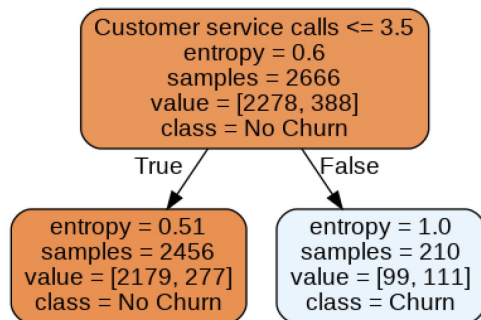


Fig: from Async

# Ensemble Exercises

[https://github.com/MIDS-W207/nteneva/tree/main/live\\_sessions\\_current/week7](https://github.com/MIDS-W207/nteneva/tree/main/live_sessions_current/week7)

KNN

# KNN

- How does it work?

# KNN

- How does it work?

## **Algorithm Pseudocode:**

- Step 1: choose the number of  $k$  (neighbors) and a distance metric.
- Step 2: Find the  $k$ -nearest neighbors of the record we want to classify.
- Step 3: Assign the class label by majority vote.

# KNN

- It's a very **different and simple** algorithm compared to the ones we will see in this course.
- Lazy learner: it doesn't learn any discriminative function from the training data, but memorizes the data instead.
- Advantages:
  - adapts easily to new data.
- Disadvantages:
  - we cannot discard training examples because no training is involved.
  - storage space can become a challenge with large datasets.
  - computationally complexity grows linearly with the size of the data.

# Implement KNN

[https://github.com/MIDS-W207/nteneva/blob/main/live\\_sessions\\_current/week7/exercise\\_knn\\_classifier.ipynb](https://github.com/MIDS-W207/nteneva/blob/main/live_sessions_current/week7/exercise_knn_classifier.ipynb)